Your work, in your own data

The twelve openstory-skills, run against your live store. The same reports as the prompt library — but every number here is yours.

> **How this was made:** each section is the real output of one of the twelve `openstory-skills` run against your live store over **the last ~30 days (arc & recall reach across all history)** — the same reports, charted from your own history. Names of collaborators are shown as roles; secrets are never included.

01 /cost · what your sessions cost

What have my agent sessions cost me?

6.46B total tokens (30d)	**~$3.2K** est. cost (Sonnet pricing)	**96%** cache hit rate	**$16.5K** saved vs. no-cache

TOKENS PER DAY

> A 96% cache hit rate turned a potential **$19.8K** bill into **$3.2K** — prompt caching saved ~$16.5K over 30 days. (Sonnet pricing: $3/M in, $0.30/M cache read, $15/M out.)

daily_token_usage (days=30)

02 /time · where your time goes

Where does my time actually go?

ACTIVITY BY HOUR (UTC, 4-HOUR BUCKETS)

1 AM UTC
peak hour (8,706 events)

27%
top-3 hours share

68.2K
total events (30d)

> The dead zone is **UTC 04–07** (near-zero) — consistent with sleep. Three active windows: UTC 12–15, 20–23, and 00–03. Not a 9-to-5 pattern — work clusters into morning, afternoon, and a late-evening push.

productivity (days=30) — hourly event density

03 /tools · what you reach for

Which tools and commands do I rely on most?

TOP TOOLS (8 SESSIONS, ~14 DAYS)

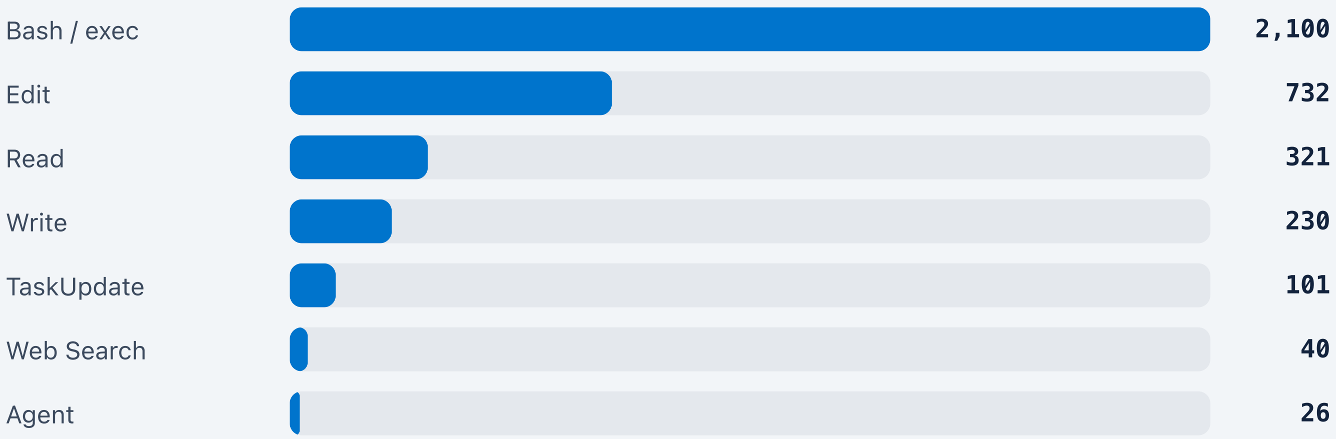

3,893 total tool calls	**2.3:1** Edit : Read ratio	**59%** Bash share

Bash dominates at 59% of visible calls — shell execution is the primary agent action. Edit outpaces Read 2.3:1: a write-heavy pattern, the agent codes more than it explores.

`list_sessions(days=14) + session_synopsis on the 8 largest sessions`

04 /coach · how you work

Honest feedback on how I work with coding agents.

SIGNAL	RATE	MEANING
Multi-day sessions	High	Sessions run 4–7 days each; context degrades and errors concentrate late in the run
Bash dominance	50% of tools	Execution-first: ops, greps, curls dominate structured reads and edits
Edit-before-Read ratio	Medium	1.85:1 — editing nearly 2x as often as reading; risk of edits without full context
Error events	Medium	23 errors in one 7-day session; the other 3 were clean — pattern is session-age correlated
Agent delegation	2% of tools	Subagents used rarely — research and exploration stay serial
Compaction hits	Low	1 compact across ~7 days; context rarely hits the hard limit

#1 pattern: **sessions run multi-day** (4–7 days each). That is where errors appear and context quality silently drifts. Fix: close and restart a session when a concrete phase ends — don't carry one thread across days of varied work.

`4 recent OpenStory sessions Jun 17–26 — 1130 tool calls total`

05 /arc · the story so far

Trace the story of how this project came to be.

2026-06-14
CI/CD deployment pipeline
Enable automatic CI/CD — 1,194 events, Bash-heavy; infra pipeline hardening

2026-06-16
Strategic reflection

Deep reflection on where the work stands — direction-setting, 2,112 events

2026-06-17

Prompt-library generators

Analysed most-common prompts → built the generators + 5-palette theme system; 2,752 events

2026-06-20

Developer-archetype research

Multi-agent adversarial fact-checking — extractor + verifier subagents; archetype charts built

2026-06-21

Architectural audit + hub connection

Full architectural assessment (1,694 events) + NATS hub connectivity in parallel

2026-06-22

Cross-agent integration debug

Getting remote-agent events into OpenStory; install troubleshooting; CLAUDE.md audit

2026-06-26

Team skill + UI wireframes

openstory:team automation run; UI review; wireframe design; codebase scan

> *Two weeks of parallel construction — CI/CD hardening, strategic reflection, prompt-library and archetype research, a deep architectural audit, cross-agent integration, and team-skill automation: OpenStory building itself while building the tooling to observe everything else.*

`list_sessions(project=OpenStory) — 200 sessions, Jun 14 → Jun 26`

06 /recap · what shipped lately

What have I worked on lately, grouped by project?

PROJECT	SESSIONS	ACTIVITY	FOCUS
OpenStory	243	15,694 events	Prompt-library generators, archetype research, UI review, NATS federation, team-skill automation
workspace	3	6,848 events	Remote agent hub sessions (OpenClaw check-ins + hub connectivity)
agent-harness	4	5,399 events	OpenActor architecture, reactive event-driven system planning
financial-research	1	3,146 events	Financial / economic research and policy analysis
openstory-deploy	1	1,194 events	CI/CD deployment pipeline automation
openstory-website	1	206 events	Website updates

> OpenStory leads at **15k+ events across 243 sessions** — including ~200 short script-layer agent sessions spawned by the prompt-library and archetype generators. Three parallel fronts: OpenStory feature work, remote agent sessions via OpenClaw (workspace), and OpenActor reactive design (agent-harness).

`project_pulse(days=14) merged by display name + list_sessions(days=14)`

07 /standup · today

Write my standup: what I did, what's blocked, what's next.

DID

- Ran `/openstory:team` skill — 583 events, ~3.3h; heavy Bash + Task automation (Jun 26)
- Ran `/openstory:scan` — codebase scan, 41 events (Jun 26)
- UI review + design discussion — assessed UI quality and direction (Jun 26)
- Wireframe exploration for a new feature — subagent design session (Jun 26)
- Two PR-review sessions — code walkthrough + review (Jun 25)

BLOCKED

- NATS leaf federation — 122-event investigation (Jun 25) into why the leaf node isn't reaching the hub; no resolution recorded

NEXT

- Resolve NATS leaf federation — continue from the Jun 25 investigation
- Iterate on UI design from the wireframe exploration + review feedback
- Follow up on openstory:team skill findings

> Last active **Jun 26**. 4 sessions / 672 events that day; 3 more (294 events) on Jun 25. Top tools: Bash, TaskUpdate, Edit — a mix of skill automation, UI review, design, and PR analysis.

`list_sessions(days=7) + session_synopsis for the top sessions on Jun 25–26`

08 /recall · find it again

Find the last time I solved this — and exactly how.

Jun 25	Make sure we are connected to the hub	OpenStory (remote)
Jun 21	Architectural assessment — touched federation config	OpenStory
older	NATS leaf bootstrap + deploy/ script	OpenStory
older	Tailnet federation spike — hub config	OpenStory
older	NATS leaf secrets + .env.federation authoring	OpenStory

```
1  curl -s http://localhost:8222/varz | python3 -c "import sys,json;
   d=json.load(sys.stdin); print('leafnodes:', d.get('leafnodes'))"
2  grep -inE 'leaf|remotes|7422|url|hub' nats.conf
3  cat .env.federation # NATS_LEAF_URL=nats://<redacted>@<redacted>:7422
4  pkill -f 'nats-server -c nats.conf' # free :4222 so the leaf node can bind it
5  nohup just up-federation > /tmp/openstory-federation-boot.log 2>&1 &
6  curl -s http://localhost:8222/leafz # out->hub / in<-hub message counts confirm the
   link
7  curl -s http://localhost:3002/api/sessions # both hosts' sessions should now appear
```

> Topic recalled: **NATS leaf federation** — connecting local OpenStory to the distributed hub over Tailscale. Last solved Jun 25, 2026. Key insight: a standalone `nats.conf` occupies `:4222` and blocks the leaf node; `up-federation` sources `.env.federation` and starts the leaf instead.

`agent_search 'federation nats leaf deploy' — 5 matches; tool_journey on the top session (Jun 21–25)`

09 /prime · pick up where you left off

Where did we leave off on this project?

LAST SESSION	Jun 26, 2026 04:06 UTC (openstory:team)
LAST GOAL	Team-day planning + UI wireframe design for a new feature
OPEN THREADS	prompt-library-generators branch; UI redesign + watch-command-center research; developer-archetype profiling
TOP TOOLS	Bash 69 · TaskUpdate 13 · Edit 10 · TaskCreate 9

Prompt-library generators `branch: tooling/prompt-library-generators`

Theme system (5 palettes) + session-usage analysis built. Several scripts still untracked — stage and PR.

Developer-archetype research `docs/research/developer-archetypes/`

Citation-traceable archetype profiling with viz scripts. Style axis refined to process-locus. Docs and scripts untracked.

UI redesign + watch command center `docs/research/ui-redesign/`

Wireframe session ran Jun 26; UI review completed same day. Feature research docs untracked.

Boot skill `.claude/skills/boot/`

New boot skill formalizing the full boot flow (Docker / native SQLite / split dev) — untracked.

Suggested next step: **stage and open a PR for the prompt-library-generators branch** — most new work (scripts/, docs/research/, the boot skill) is still untracked.

`project_context for OpenStory (last 5 sessions Jun 25–26) + git status on the active branch`

10 /watch · live work

Watch the work happening right now and summarise it.

04:06	Bash	Check UI build log — dist dir present after background npm run build
03:47	Read	Read rs/cli/src/main.rs — inspect --static-dir serve flag
03:46	Bash	Search build_router + static serve wiring across rs/ crates
03:46	Bash	Launch UI production build in background, check vite proxy config
03:45	Bash	Inspect CLI serve subcommand args and vite.config.ts proxy settings
03:41	AskUserQuestion	Ask: serve new WatchBanner build via --static-dir or keep Vite proxy?
03:41	Bash	Verify open-story serve + Vite processes on :3002 / :5173
03:40	Bash	Test /api/sessions response count + WS upgrade path live

Static snapshot of session **d58d3d03** (started 00:46 UTC, Jun 26). In the live deck this feed streams in real time from the `/ws` WebSocket. This session built the watch-focus feature — `POST /api/watch/{id}` + `WatchBanner` UI — the same mechanism `/watch` uses to aim the UI at a session.

`tool_journey for most recent session d58d3d03 (last_event 2026-06-26T04:06 UTC, 583 events)`

11 /team · who's working

Who on my team is active, and on what?

You · local 29 sessions

OpenStory core work: agent-directed watch-focus endpoint (POST /api/watch + WatchBanner UI), PR reviews, CI/CD deploy automation, prompt-library generators, openstory:team/scan/coach skills, CLAUDE.md audit, NATS federation debugging

You · remote 9 sessions

Remote Linux box: NATS leaf node federation, OpenActor/agent-harness development, cross-machine OpenStory connectivity, economic research one-off

OpenClaw agent `3 sessions`

Automated workspace reads (HEARTBEAT.md context pattern on startup); long-running session spanning months with periodic activation from the VPS

OpenActor (automated) `1 session`

Single automated benchmark run (coding QA task, short-lived)

> This is a **single-operator store** — all sessions belong to you across two machines plus two agents. `/team` lights up with real multi-user data when a teammate's NATS leaf node federates in: their sessions arrive with a distinct host prefix and appear as a separate card here automatically.

`list_sessions days=7 (56 sessions); project_id path prefixes identify 4 distinct host/agent segments`

12 /scan · before you share

Is there anything sensitive in my history?

CATEGORY	HITS	STATUS
`sk-` API key prefix	0	Low
`ghp_` GitHub token prefix	0	Low
`AKIA` AWS key prefix	0	Low
`BEGIN PRIVATE KEY`	0	Low
`xoxb` / `xoxp` Slack token	0	Low
`password` keyword	10	Medium
`Bearer` / `api_token`	0	Low

PATTERNS SCANNED `sk-` `ghp_` `AKIA` `BEGIN PRIVATE KEY` `xoxb` `xoxp` `password` `Bearer` `api_token`

> All high-signal key prefixes are **zero** — confirmed clean. The 10 `password` hits are session-transcript context only: SSH login prompts from a server setup and dummy test-fixture credentials in a NATS multi-account test config — no real credentials. `Bearer`/`api_token` hits are Rust middleware source snippets read during code review, not live token values.

`FTS5 search across live event store for 7 sensitive-pattern categories; baseline confirmed by dedicated /scan session (80b66f5f) which ran grep against the transcript tree`